



# Problem Statement

- To predict the best fit model for predicting fraudulent activities in credit card dataset.

# Data Description

- The dataset consists of 24 columns:
- Independent Variable: Limit balance Sex, Education, Marriage, Age, Pay and Bill
- Dependent Variable: Default Payment (y)
- There are no Missing Variable

```
dt.head()
```

|   | LIMIT_BAL | SEX | EDUCATION | MARRIAGE | AGE | PAY_0 | PAY_2 | PAY_3 | PAY_4 | PAY_5 | ... | BILL_AMT4 | BILL_AMT5 | BILL_AMT6 | PAY_AMT1 | PAY_AMT2 | PAY_AMT3 | PAY_AMT4 | PAY_AMT5 | PAY_AMT6 | y |
|---|-----------|-----|-----------|----------|-----|-------|-------|-------|-------|-------|-----|-----------|-----------|-----------|----------|----------|----------|----------|----------|----------|---|
| 0 | 20000     | 2   | 2         | 1        | 24  | 2     | 2     | -1    | -1    | -2    | ... | 0         | 0         | 0         | 0        | 689      | 0        | 0        | 0        | 0        | 1 |
| 1 | 120000    | 2   | 2         | 2        | 26  | -1    | 2     | 0     | 0     | 0     | ... | 3272      | 3455      | 3261      | 0        | 1000     | 1000     | 1000     | 0        | 2000     | 1 |
| 2 | 90000     | 2   | 2         | 2        | 34  | 0     | 0     | 0     | 0     | 0     | ... | 14331     | 14948     | 15549     | 1518     | 1500     | 1000     | 1000     | 1000     | 5000     | 0 |
| 3 | 50000     | 2   | 2         | 1        | 37  | 0     | 0     | 0     | 0     | 0     | ... | 28314     | 28959     | 29547     | 2000     | 2019     | 1200     | 1100     | 1069     | 1000     | 0 |
| 4 | 50000     | 1   | 2         | 1        | 57  | -1    | 0     | -1    | 0     | 0     | ... | 20940     | 19146     | 19131     | 2000     | 36681    | 10000    | 9000     | 689      | 679      | 0 |

5 rows × 24 columns

| No        | Attribute Name | Description   |
|-----------|----------------|---|
| <b>1</b>  | ID             | User ID   |
| <b>2</b>  | Limit_Bal      | Amount of the given credit (NT dollar)              |
| <b>3</b>  | Sex            | Gender(1=male,2=female)                             |
| <b>4</b>  | Education      | Education (1=graduate school,2=University,3=others) |
| <b>5</b>  | Marriage       | Marital status(1=married,2=unmarried)               |
| <b>6</b>  | Age            | Age(year)   |
| <b>7</b>  | Pay_0          | the repayment status in September, 2005             |
| <b>8</b>  | Pay_2          | the repayment status in August, 2005                |
| <b>9</b>  | Pay_3          | the repayment status in July, 2005                  |
| <b>10</b> | Pay_4          | the repayment status in June, 2005                  |
| <b>11</b> | Pay_5          | the repayment status in May, 2005                   |
| <b>12</b> | Pay_6          | the repayment status in April, 2005                 |
| <b>13</b> | Bill_Amt1      | amount of bill statement in September, 2005         |
| <b>14</b> | Bill_Amt2      | amount of bill statement in August, 2005            |
| <b>15</b> | Bill_Amt3      | amount of bill statement in July, 2005              |

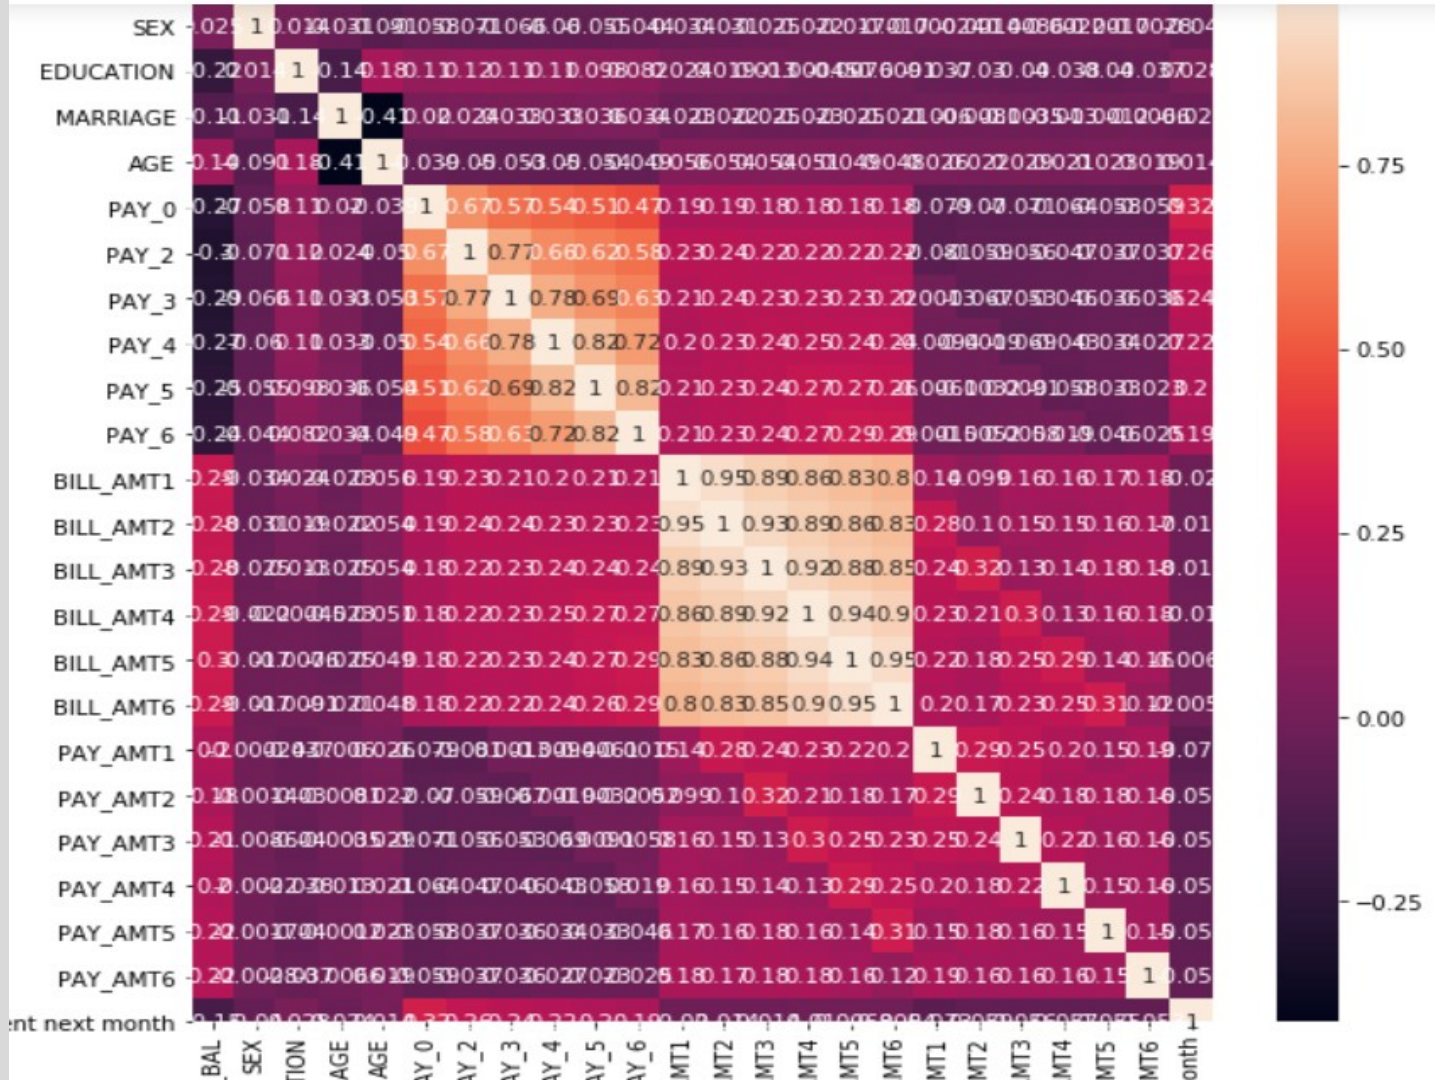
|           |           |   |
|-----------|-----------|---|
| <b>16</b> | Bill_Amt4 | amount of bill statement in June, 2005  |
| <b>17</b> | Bill_Amt5 | amount of bill statement in May, 2005   |
| <b>18</b> | Bill_Amt6 | amount of bill statement in April, 2005 |
| <b>19</b> | Pay_Amt1  | amount paid in April, 2005              |
| <b>20</b> | Pay_Amt2  | amount paid in May, 2005                |
| <b>21</b> | Pay_Amt3  | amount paid in June, 2005               |
| <b>22</b> | Pay_Amt4  | amount paid in July, 2005               |
| <b>23</b> | Pay_Amt5  | amount paid in August, 2005             |
| <b>24</b> | Pay_Amt6  | amount paid in September, 2005          |
| <b>25</b> | Default   | Amount to be paid next month            |

## CORRELATION

Heat map is plotted to check how well the variables are correlated with each other.

We can see that Education and Repayment status represents strong negative correlation with Amount of credit Given (LIMIT\_BAL)

Each Variables of Billing statement:  
BILL\_AMT1,  
BILL\_AMT6 are positively correlated among each other.



# Scaling - Standard Scaler

|   | LIMIT_BAL | SEX       | EDUCATION | MARRIAGE  | AGE       | PAY_0     | PAY_2    | PAY_3     | PAY_4     | PAY_5     | ... | BILL_AMT4 | BILL_AMT5 | BILL_AMT6 |
|---|-----------|-----------|-----------|-----------|-----------|-----------|----------|-----------|-----------|-----------|-----|-----------|-----------|-----------|
| 0 | -1.136720 | 0.810161  | 0.185828  | -1.057295 | -1.246020 | 1.794564  | 1.782348 | -0.696663 | -0.666599 | -1.530046 | ... | -0.672497 | -0.663059 | -0.652724 |
| 1 | -0.365981 | 0.810161  | 0.185828  | 0.858557  | -1.029047 | -0.874991 | 1.782348 | 0.138865  | 0.188746  | 0.234917  | ... | -0.621636 | -0.606229 | -0.597966 |
| 2 | -0.597202 | 0.810161  | 0.185828  | 0.858557  | -0.161156 | 0.014861  | 0.111736 | 0.138865  | 0.188746  | 0.234917  | ... | -0.449730 | -0.417188 | -0.391630 |
| 3 | -0.905498 | 0.810161  | 0.185828  | -1.057295 | 0.164303  | 0.014861  | 0.111736 | 0.138865  | 0.188746  | 0.234917  | ... | -0.232373 | -0.186729 | -0.156579 |
| 4 | -0.905498 | -1.234323 | 0.185828  | -1.057295 | 2.334029  | -0.874991 | 0.111736 | -0.696663 | 0.188746  | 0.234917  | ... | -0.346997 | -0.348137 | -0.331482 |

5 rows × 24 columns

The Standard Scaler transforms the data such that each value will have a mean 0 and Standard Deviation 1.

| PAY_AMT1  | PAY_AMT2  | PAY_AMT3  | PAY_AMT4  | PAY_AMT5  | PAY_AMT6  | default payment next month |
|-----------|-----------|-----------|-----------|-----------|-----------|----------------------------|
| -0.341942 | -0.227086 | -0.296801 | -0.308063 | -0.314136 | -0.293382 | 1                          |
| -0.341942 | -0.213588 | -0.240005 | -0.244230 | -0.314136 | -0.180878 | 1                          |
| -0.250292 | -0.191887 | -0.240005 | -0.244230 | -0.248683 | -0.012122 | 0                          |
| -0.221191 | -0.169361 | -0.228645 | -0.237846 | -0.244166 | -0.237130 | 0                          |
| -0.221191 | 1.335034  | 0.271165  | 0.266434  | -0.269039 | -0.255187 | 0                          |

# Outlier(After Scaling)

- Outliers are extreme values that deviate from other observations on data
- Z score - Since 99.73 data points lie within  $\pm 3$  Standard Deviation. if Z is greater than 3 then that value is considered as Outliers

```
#Removing outlier - z_score
from scipy import stats
z = np.abs(stats.zscore(d1))
print(z)
```

```
[[1.13672015  0.81016074  0.18582826 ... 0.31413612  0.29338206  1.87637834]
 [0.3659805   0.81016074  0.18582826 ... 0.31413612  0.18087821  1.87637834]
 [0.59720239  0.81016074  0.18582826 ... 0.24868274  0.01212243  0.53294156]
 ...
 [1.05964618  1.23432296  0.18582826 ... 0.18322937  0.11900109  1.87637834]
 [0.67427636  1.23432296  1.45111372 ... 3.15253642  0.19190359  1.87637834]
 [0.90549825  1.23432296  0.18582826 ... 0.24868274  0.23713013  1.87637834]]
```

```
#Removing outliers
threshold = 3
d1.shape
```

```
(30000, 24)
```

```
#creating new dataframe after removing outliers
d2 = d1[(z<3).all(axis=1)]
d2.shape
```

```
(26429, 24)
```

# Outlier Detection(Without Scaling)

```
from scipy import stats
z = np.abs(stats.zscore(df[['X1', 'X5', 'X12', 'X13', 'X14', 'X15', 'X16', 'X17', 'X18', 'X19', 'X20', 'X21', 'X22', 'X23']]))
print(z)
```

```
[[1.13584859 1.24613545 0.64373456 ... 0.30844347 0.31452581 0.29374281]
 [0.36525716 1.02923602 0.66044213 ... 0.24467821 0.31452581 0.1813593 ]
 [0.59643459 0.16163827 0.30000105 ... 0.24467821 0.24914153 0.01278404]
 ...
 [1.05878945 0.16371088 0.64845774 ... 0.04062937 0.18375725 0.11954838]
 [0.67349373 0.59750975 0.71916972 ... 0.18563158 3.14848722 0.19237289]
 [0.90467116 1.13975835 0.0463337 ... 0.24467821 0.24914153 0.23755106]]
```

```
threshold = 3
print(np.where(z > 3))
```

```
(array([ 6, 6, 6, ..., 29923, 29928, 29928], dtype=int64), array([ 2, 3, 4, ..., 9, 8, 12], dtype=int64))
```

```
df.shape
```

```
(29930, 24)
```

```
df_new = df[(z < 3).all(axis=1)]
```

```
df_new.shape
```

```
(27007, 24)
```




# Smote - Imbalanced dataset

(Synthetic Minority Oversampling Technique)

```
{1: 5935, 0: 20494}  
{0: 20493, 1: 16395}
```

```
#To check whether data is balanced or not  
y1.hist() #seen imbalanced data
```

```
<AxesSubplot:>
```





# Variable exclusion - LOGIT I

- To check for insignificant variables and remove them

This is the **final output** after removing the necessary insignificant data and thus getting all the variables that are significant.

Since the *LLR p-value is lesser than 0.05 and all variables are significant*, the model is deployable and these variables could be used in implementation of further techniques.

|           |         |       |         |       |        |        |
|-----------|---------|-------|---------|-------|--------|--------|
| BILL_AMT1 | -0.6364 | 0.050 | -12.611 | 0.000 | -0.735 | -0.537 |
| BILL_AMT3 | 0.2276  | 0.065 | 3.487   | 0.000 | 0.100  | 0.356  |
| BILL_AMT6 | 0.1702  | 0.040 | 4.286   | 0.000 | 0.092  | 0.248  |
| PAY_AMT1  | -0.1001 | 0.039 | -2.557  | 0.011 | -0.177 | -0.023 |
| PAY_AMT2  | -0.2517 | 0.051 | -4.977  | 0.000 | -0.351 | -0.153 |
| PAY_AMT6  | 0.1975  | 0.039 | 5.030   | 0.000 | 0.121  | 0.274  |

## Logit Regression Results

|                  |                            |                   |         |
|------------------|----------------------------|-------------------|---------|
| Dep. Variable:   | default payment next month | No. Observations: | 36888   |
| Model:           | Logit                      | Df Residuals:     | 36872   |
| Method:          | MLE                        | Df Model:         | 15      |
| Date:            | Tue, 13 Oct 2020           | Pseudo R-squ.:    | 0.08817 |
| Time:            | 23:18:11                   | Log-Likelihood:   | -23106. |
| converged:       | True                       | LL-Null:          | -25341. |
| Covariance Type: | nonrobust                  | LLR p-value:      | 0.000   |

|           | coef    | std err | z      | P> z  | [0.025 | 0.975] |
|-----------|---------|---------|--------|-------|--------|--------|
| LIMIT_BAL | -0.0743 | 0.015   | -4.879 | 0.000 | -0.104 | -0.044 |
| SEX       | -0.0632 | 0.011   | -5.613 | 0.000 | -0.085 | -0.041 |
| EDUCATION | -0.0288 | 0.014   | -2.121 | 0.034 | -0.055 | -0.002 |
| MARRIAGE  | -0.0747 | 0.012   | -6.019 | 0.000 | -0.099 | -0.050 |
| AGE       | 0.0595  | 0.013   | 4.556  | 0.000 | 0.034  | 0.085  |
| PAY_0     | 0.5741  | 0.015   | 38.184 | 0.000 | 0.545  | 0.604  |
| PAY_2     | 0.1737  | 0.019   | 9.138  | 0.000 | 0.136  | 0.211  |
| PAY_3     | 0.0476  | 0.021   | 2.261  | 0.024 | 0.006  | 0.089  |
| PAY_4     | 0.0420  | 0.020   | 2.091  | 0.037 | 0.003  | 0.081  |
| PAY_6     | -0.0378 | 0.017   | -2.232 | 0.026 | -0.071 | -0.005 |

# Variable exclusion - Principal Component Analysis II

```
#PCA for feature selection and reduce dimensionality  
#Here instead of mentioning n_components, we will directly mention the variance we would like PCA to capture  
from sklearn.decomposition import PCA  
pca = PCA(0.85)  
pca.fit(x_train)  
PCA(n_components=0.85, random_state=None, whiten = False)
```

```
PCA(n_components=0.85)
```

```
#Checking for number of components that would provide with 85% variance  
pca.n_components_
```

```
8
```

```
#Just to cross check whether 8 components would be enough to explain major variance in the data  
pca = PCA(n_components=8)  
fit = pca.fit(x)  
print("Explained Variance: %s" % fit.explained_variance_ratio_)  
print(fit.components_)
```

```
Explained Variance: [0.36959408 0.15302819 0.10437215 0.0724056  0.05658796 0.04980261  
0.03840171 0.03293198]
```

- Applied PCA, removed the variables that were insignificant according to the given here. (Having least variance among each other)
- After transforming the dataset, different techniques were applied using the rest of the significant variables.

# Techniques - Logistic Regression (After LOGIT)

- Applied on the significant variables derived above.
- **Model Assumptions:**
  - Target variable is binary
  - Predictive features are interval (continuous) or categorical
  - Features are independent of one another
  - Sample size is adequate
- **Accuracy : 0.694**
- **F1 - Score : 0.624**

```
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
cm_lr = confusion_matrix(y_test, y_pred_test)
acc_lr = accuracy_score(y_test, y_pred_test)
pre_lr = precision_score(y_test, y_pred_test)
recall_lr = recall_score(y_test, y_pred_test)
f1_lr = f1_score(y_test, y_pred_test)
```

```
print('Confusion Matrix : \n',cm_lr)
print('Accuracy : ',acc_lr)
print('Precision score', pre_lr)
print('Recall score',recall_lr)
print('F1 score',f1_lr)
```

```
Confusion Matrix :
[[4051 1086]
 [1738 2347]]
Accuracy : 0.6937757536326177
Precision score 0.683658607631809
Recall score 0.5745410036719706
F1 score 0.624368183027401
```

# Techniques - Logistic Regression (After PCA)

- **Accuracy : 0.6025**
- **F1 - Score : 0.3121**

```
from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score
cm_lr = confusion_matrix(y_test, y_pred_test)
acc_lr = accuracy_score(y_test, y_pred_test)
pre_lr = precision_score(y_test, y_pred_test)
recall_lr = recall_score(y_test, y_pred_test)
f1_lr = f1_score(y_test, y_pred_test)

print('Confusion Matrix : \n',cm_lr)
print('Accuracy : ',acc_lr)
print('Precision score', pre_lr)
print('Recall score',recall_lr)
print('F1 score',f1_lr)
```

Confusion Matrix :

```
[[4801  445]
 [3279  845]]
```

Accuracy : 0.6025613660618997

Precision score 0.6550387596899225

Recall score 0.20489815712900097

F1 score 0.31215367565570745

# Techniques - Decision Tree (After LOGIT)

- Similar approach in case of decision tree.
- **Accuracy : 0.746**
- **F1 - score : 0.718**
- Though better than that in the case of logistic regression, but will try to impute the improvised version of the same
  - Random Forest to check if the accuracy is increasing or not

Confusion Matrix :

```
[[3890 1247]  
 [1098 2987]]
```

Accuracy : 0.7457167642593797

Precision score : 0.7054794520547946

Recall score : 0.7312117503059975

F1 score : 0.7181151580718836

# Techniques - Decision Tree (After PCA)

- Similar approach in case of decision tree.
- **Accuracy : 0.7239**
- **F1 - score : 0.6911**

Confusion Matrix :

```
[[3889 1357]
```

```
[1230 2894]]
```

Accuracy : 0.723906083244397

Precision score : 0.6807809927075982

Recall score : 0.7017458777885548

F1 score : 0.6911044776119403

# Techniques - Random Forest Classifier(After LOGIT)

- Being improvised version of Decision Tree, it resulted in deriving **higher** accuracy and F1 score.
- **Accuracy : 0.829**
- **F1 Score : 0.797**

Confusion Matrix :

```
[[4554  583]
```

```
[ 994 3091]]
```

Accuracy : 0.8289958794187812

Precision score : 0.8413173652694611

Recall score : 0.7566707466340269

F1 score : 0.7967521587833484

# Techniques - Random Forest Classifier(After PCA)

- Being improvised version of Decision Tree, it resulted in deriving **higher** accuracy and F1 score.
- **Accuracy : 0.822**
- **F1 Score : 0.7862**

Confusion Matrix :

```
[[4654  592]  
 [1069 3055]]
```

Accuracy : 0.8227321237993597

Precision score : 0.8376748012064711

Recall score : 0.7407856450048497

F1 score : 0.7862565950328143



# Techniques - Support Vector Machine (SVM)

AFTER LOGIT

- Another Powerful Machine Learning Algorithm which is used for both Classification and Regression.
- **Accuracy : 0.726**
- **F1 Score : 0.647**

Confusion Matrix :

```
[[4384  753]
 [1771 2314]]
```

Accuracy : 0.7263066579917589

Precision score : 0.7544832083469188

Recall score : 0.5664626682986537

F1 score : 0.6470917225950783

# Techniques - Support Vector Machine (SVM)

## AFTER PCA

- Since the target variable is **binary in nature**, we tried to implement SVM as it is a binary classifier.
- **Accuracy : 0.63299**
- **F1 Score : 0.5545**

Confusion Matrix :

```
[[3795 1451]
```

```
[1985 2139]]
```

Accuracy : 0.6332977588046959

Precision score : 0.5958217270194986


Recall score : 0.5186711930164889

F1 score : 0.5545760954109411

# Techniques - K-Nearest Neighbours (KNN): After LOGIT

Best K is : 3 | Cross validation Accuracy : 0.7653067306790259

```
#Plotting CV accuracy curve for clarity  
plt.plot(range(3,12),accList) #Scores list  
plt.xlabel('K')  
plt.ylabel('CV Accuracy')  
plt.show()
```



Confusion Matrix :

```
[[3598 1539]  
 [ 594 3491]]
```

Accuracy : 0.7687052700065062

Precision score : 0.6940357852882704


Recall score : 0.8545899632802938

F1 score : 0.7659901261656611

- Here, we have given a **range from 3-12** as **k**, in order **to estimate the best k** among the above which can derive better accuracy when compared.
- In this model, **k = 3** turned out to be the best K, having the **highest CV accuracy** than others.
- The CV Accuracy plot clearly shows that as the number of K increases the CV accuracy falls.

# Techniques - K-Nearest Neighbours (KNN): After PCA

```
#Plotting CV accuracy curve for clarity
plt.plot(range(3,12),accList) #Scores list
plt.xlabel('K')
plt.ylabel('CV Accuracy')
plt.show()
```



Confusion Matrix :

```
[[3527 1719]
```

```
[ 738 3386]]
```

Accuracy : 0.7377801494130203

Precision score : 0.6632713026444662

Recall score : 0.8210475266731329

F1 score : 0.733773973344891

- Here, we have given a **range from 3-12** as **k**, in order **to estimate the best k** among the above which can derive better accuracy when compared.
- In this model, **k =3** turned out to be the best K, having the **highest CV accuracy** than others.
- The CV Accuracy plot clearly shows that as the number of K increases the CV accuracy falls.

# Conclusion - Best Model

## LOGIT

| Techniques          | Accuracy | F1_Score |
|---------------------|----------|----------|
| Logistic Regression | 0.694    | 0.624    |
| Decision Tree       | 0.742    | 0.716    |
| Random Forest       | 0.831    | 0.8      |
| SVM                 | 0.725    | 0.645    |
| KNN                 | 0.781    | 0.777    |

## PCA

| Techniques          | Accuracy | F1_Score |
|---------------------|----------|----------|
| Logistic Regression | 0.603    | 0.312    |
| Decision Tree       | 0.724    | 0.691    |
| Random Forest       | 0.822    | 0.786    |
| SVM                 | 0.633    | 0.555    |
| KNN                 | 0.738    | 0.733    |



# Q&A

## 1. Which other scalar techniques did you use, why did you only consider standard scalar?

Ans. We used Min-Max scalar and Robust scalar apart from standard scalar. While applying Min-Max scalar we got an error that the smote data was not able to fit in the model. Hence, we used Standard scaler for standardizing smote data.

We also used Robust scaler after removing outliers using Isolation Forest.

## 2. Explain the correlation matrix

Ans. From the heat map we could make out that is Limit\_Bal, which means the amount limit the credit card gets to spend on credit is highly negatively correlated with the default. Which means that higher the credit limit lower the chance of the customer to default. This is because your credit card limit increases on your ability to payback and the account balance you have so if you are a safer customer, then only you have higher credit limit and the chance of defaulting is low. Limit balance is the main factor affecting the default of the customer.



# Business Applications of Fraud Detection Models in the Finance World

## 1. **Predictive Modelling**

Use Historical Behavioural information of known fraud to identify suspicious behaviour similar to previous fraud patterns.

## 2. **Anomaly Detection**

It helps to identify data points and observations that deviate from a dataset's normal behaviour.

## 3. **Natural Language Processing**

NLP with Machine Learning is a win win combination used to detect fraud. It can be used for :

- (a) Transaction Narration Analysis
- (b) Company News and Sentiment Analysis
- (c) Email Monitoring
- (d) Financial Statement Analysis



**THANK YOU**